

Seminar Report

On

KrishiSaarthi: An Integrated AI Platform for Smallholder Farm Management

Submitted by

*Rohan
Agarwal*

12216367

**Bachelor of
Technology**

IN

(Computer Science and Engineering)

Under the Supervision of

*Ojasvee
Kaneria*

*(Assistant
Professor)*



**LOVELY PROFESSIONAL
UNIVERSITY PUNJAB**

(April, 2026)

DECLARATION

I hereby declare that the seminar report titled **“KrishiSaarthi: An Integrated AI Platform for Smallholder Farm Management”** submitted in partial fulfilment of the requirements for the award of the degree of Bachelor of Technology in Computer Science and Engineering is a record of my own work carried out during the academic session **2025-2026**.

I further declare that this report has not been submitted, either in part or in full, to any other institution or university for the award of any degree or diploma.

I confirm that the content of this report is original and prepared by me. Any references used have been duly acknowledged. I also declare that the use of Artificial Intelligence (AI) tools, if any, has been minimal and the AI-generated content in this report is **less than 10%**, ensuring that the majority of the work reflects my own understanding and effort.

I take full responsibility for the authenticity and originality of the work presented in this report.

Name of the Student: Rohan Agarwal

Registration Number: 12216367

Course: B.Tech CSE

Signature of the Student: Rohan

Date: 10-04-2026

KrishiSaarthi: An Integrated AI Platform for Smallholder Farm Management

Final Capstone Project Report

Academic Year: 2025–26

1. [Chapter 1: Introduction]
2. [Chapter 2: Literature Review]
3. [Chapter 3: Conceptual Study / System Architecture & Design]
4. [Chapter 4: Results and Discussion]
5. [Chapter 5: Conclusion and Future Scope]
6. [References]

Chapter-1

Introduction

1.1 Title of the Project

KrishiSaarthi: An Integrated AI Platform for Smallholder Farm Management

1.2 Background and Importance of the Topic

According to the Food and Agriculture Organization (FAO), smallholder farms—defined as farms under two hectares—account for approximately 84% of the world's farms and contribute 35% of the global food supply. Despite their critical importance to food security, these farmers face a confluence of challenges: crop failures due to pest infestations and disease outbreaks, price volatility in agricultural markets, mounting pest pressure from changing ecosystems, and the aggravating effects of climate change on traditional farming practices.

Precision agriculture has emerged as a transformative paradigm that leverages satellite data, cloud computing, machine learning, and Internet of Things (IoT) technologies to deliver field-level decision support. However, most existing solutions are designed for large commercial farms and require prohibitive licensing fees, specialized hardware, or advanced technical literacy—effectively excluding the vast majority of smallholder farmers from the benefits of modern agricultural technology.

The specific gaps in existing smallholder tooling that motivate this work include:

7. **Manual Crop Health Monitoring** — Crop health monitoring remains predominantly manual and depends on physical inspection, which is time-consuming, subjective, and error-prone.
8. **Expert-Dependent Disease Identification** — Disease identification demands taxonomic expertise that is largely unavailable in rural communities, leading to delayed or incorrect diagnoses.
9. **Informal Financial Record-Keeping** — Informal financial management prevents access to institutional credit and obscures true profitability.
10. **Asymmetric Market Pricing** — Information asymmetry in market pricing disadvantages smallholder sellers, who lack access to real-time commodity data.

11. **Untapped Carbon Credit Revenue** — Sustainable Alternate Wetting and Drying (AWD) irrigation practices generate quantifiable carbon credits, yet farmers have no mechanism to document and monetize this environmental benefit.
12. **Language and Literacy Barriers** — English-only text platforms exclude multilingual and low-literacy users from accessing agricultural advisory services.

KrishiSaarthi (Sanskrit: "Agriculture Companion") is designed to address all six of these dimensions within a single, production-ready, open-source platform.

1.3 Objectives of the Project

Three research objectives guide this work:

- **RO1** — Design and evaluate KrishiSaarthi as a production-grade platform with a complete REST API covering all identified smallholder needs.
- **RO2** — Document and empirically evaluate the MobileNetV2-based Convolutional Neural Network (CNN), two-layer Long Short-Term Memory (LSTM) risk forecasting model, and composite health scoring model against published benchmarks.
- **RO3** — Verify platform coverage across all six identified smallholder needs and benchmark against leading commercial tools and recent peer-reviewed literature.

1.4 Brief Overview of the Approach and Methodology

KrishiSaarthi is a full-stack web application comprising four major functional layers:

13. **Satellite Vegetation Analysis** — Real-time computation of vegetation indices (NDVI, EVI, SAVI, NDWI) via the Google Earth Engine Python API over user-defined field polygons using cloud-free Sentinel-2 L2A imagery.
14. **38-Class MobileNetV2 Crop Disease Detection** — A deep learning CNN classifier trained on the PlantVillage dataset, preceded by a Gemini Vision semantic pre-validation stage to eliminate non-plant image false positives.
15. **LSTM Temporal Risk Prediction** — A two-layer LSTM neural network trained on real-world ERA5 climate reanalysis data covering 25 Indian agricultural districts (January 2023 – December 2024), producing daily binary risk forecasts from 10-day input sequences.

16. **Conversational Advisory** — An LLM-powered chat interface using Google Gemini 2.5 Flash, providing context-aware agricultural advice with multilingual support (English, Hindi, Punjabi, Malayalam) and voice input capability.

The platform is built using Django 5.1 with Django REST Framework on the backend (45 REST endpoints across 8 modules) and React 18 with TypeScript on the frontend (79 components, 21 dashboard modules). Both development and production deployment configurations are provided via Docker Compose, with Prometheus/Grafana observability and GitHub Actions CI/CD.

1.5 Scope of the Project

The scope of this project encompasses the complete design, development, deployment, and empirical evaluation of the KrishiSaarthi platform. Specifically, the project covers:

17. **Geographic Scope** — The platform is designed and validated for Indian agricultural conditions, with ERA5 climate data covering 25 representative agricultural districts across diverse agro-climatic zones including arid (Rajasthan, Gujarat), semi-arid (Haryana, Maharashtra), humid tropical (Kerala, Tamil Nadu), and sub-tropical (Punjab, Uttar Pradesh) regions.
18. **Crop Coverage** — The MobileNetV2 disease classifier supports 38 classes spanning 14 crop species available in the PlantVillage dataset, including tomato (10 classes), grape (4 classes), apple (4 classes), corn (4 classes), potato (3 classes), pepper (2 classes), strawberry (2 classes), cherry (2 classes), peach (2 classes), citrus (1 class), blueberry (1 class), raspberry (1 class), soybean (1 class), and squash (1 class).
19. **Functional Scope** — The platform delivers end-to-end farm management capabilities across six functional verticals: satellite-based crop monitoring, AI-powered disease detection, temporal risk forecasting, financial management with P&L analysis, carbon credit estimation, and multilingual conversational advisory.
20. **Technical Scope** — The system is implemented as a production-ready containerized application with complete CI/CD pipeline, monitoring infrastructure, and documentation. All machine learning models are designed for CPU-only inference to ensure deployability on standard commodity hardware without GPU requirements.

Exclusions: The project does not include real-world field trials with actual farmers (identified as future work), IoT sensor integration, drone-based imagery analysis, or

blockchain-based supply chain tracking. The disease detection model is evaluated exclusively on the PlantVillage laboratory dataset; field-level validation remains a planned future activity.

1.6 Significance and Expected Impact

The significance of KrishiSaarthi lies in its holistic approach to addressing the multifaceted challenges faced by smallholder farmers through a single, integrated, and accessible platform. The key contributions and expected impacts are:

21. **Democratization of Precision Agriculture** — By providing an open-source, MIT-licensed platform that requires no specialized hardware (GPU-free inference, standard web browser access), KrishiSaarthi removes the primary barriers—cost, complexity, and technical literacy—that have historically excluded smallholder farmers from precision agriculture technologies. The estimated total cost of ownership is reduced by over 90% compared to commercial alternatives such as Climate FieldView and CropX, which require annual licensing fees of USD 500–2,000 per farm.
22. **Novel Carbon Economics Integration** — The IPCC Tier 2 AWD carbon credit estimation module introduces a previously unexplored revenue pathway for smallholder rice farmers. With an estimated potential of ₹1,400–₹10,800 per AWD cycle depending on field size, this feature can provide meaningful supplemental income while simultaneously incentivizing environmentally sustainable irrigation practices that reduce methane emissions by up to 48%.
23. **Bridging the Language Divide** — The multilingual support (English, Hindi, Punjabi, Malayalam) with voice input capability addresses the critical accessibility gap identified in existing platforms, which are predominantly English-only. This is particularly significant given that only 10.6% of India's population speaks English as a first or second language (Census 2011), and approximately 25% of smallholder farmers have limited formal literacy—making voice-based interaction an essential accessibility feature.
24. **Advancing the State of the Art** — The two-stage LLM-augmented inference pipeline (Gemini Vision pre-filtering followed by MobileNetV2 CNN classification) represents a novel architectural contribution that addresses the well-documented lab-to-field accuracy degradation problem in crop disease detection. This contribution is entirely absent from all 17 reviewed works in the literature survey.

25. Reproducibility and Community Building — As a fully open-source project with comprehensive documentation, Docker-based deployment, and a complete REST API with 45 endpoints, KrishiSaarthi serves as both a deployable tool for immediate use and a research platform for the broader agricultural AI community. The MIT license permits unrestricted modification and redistribution, enabling NGOs, agricultural extension services, and government agencies to adapt the platform to local requirements.

1.7 Organization of the Report

This report is organized into five chapters following a logical progression from context and motivation through design, evaluation, and conclusions:

- **Chapter-1: Introduction** — Presents the background and motivation, defines the research objectives, describes the project scope, and discusses the expected significance and impact of the work.
- **Chapter-2: Literature Review** — Provides a comprehensive review of 17 peer-reviewed works across five thematic clusters (disease detection, yield prediction, temporal forecasting, federated learning, and LLM integration), identifies convergent themes, and maps the specific research gaps that KrishiSaarthi addresses.
- **Chapter-3: Conceptual Study / System Architecture & Design** — Describes the core theoretical concepts underlying the platform, presents the four-tier system architecture, details the technology stack, documents the machine learning pipeline (CNN, LSTM, health score), and provides workflow diagrams for key processes.
- **Chapter-4: Results and Discussion** — Reports empirical evaluation results for all ML models including per-class CNN analysis with 99.03% top-1 accuracy, LSTM risk forecasting with 97.56% accuracy and AUC-ROC of 0.9983, explainability analysis using Grad-CAM++ and SHAP, AWD detection and carbon credit estimation results, composite health score validation, and comparative positioning against commercial platforms and peer-reviewed literature.
- **Chapter-5: Conclusion and Future Scope** — Summarizes the work, documents major learning outcomes, draws five key conclusions, and identifies six future research directions with an implementation roadmap.

Chapter-2

Literature Review

2.1 Overview of the Review

A comprehensive review of 17 peer-reviewed papers was conducted, covering five thematic clusters in precision agriculture: disease detection, yield prediction, temporal forecasting, federated learning, and LLM integration. Nine representative works from this broader review are summarized below.

2.2 Summary of Reviewed Works

Ref.	Authors & Year	Key Results	Gap / Future Scope
[1]	Manzoor (2024)	Accurate disease and yield detection using ML techniques	Scalability; regional validation
[2]	Sudha & Loret (2026)	CNNs achieved 94.5%, LSTMs 92.7% accuracy	Hybrid ML+XAI; blockchain integration
[3]	Mazumder et al. (2024)	99.99% cassava, 94.94% wheat disease detection	Multi-species + XAI layers
[4]	Alkanan & Gulzar (2024)	MobileNetV2: 96.27%; precision 0.975	Diverse multi-crop datasets
[5]	Meghraoui et al. (2024)	CNN-LSTM dominates; satellite ~50%	XAI + high resolution
[6]	Wang et al. (2024)	R ² up to 0.98 for corn yield prediction	Explainability; multimodal fusion
[7]	Jabed & Murad (2024)	CNN-LSTM best for regional prediction	XAI + satellite-IoT integration
[8]	Hiremani et al. (2025)	Federated Learning ≥97%; privacy maintained	FL+XAI for edge devices
[9]	Shaikh et al. (2025)	LLMs excel in zero-shot agriculture tasks	Ag-specific foundation models

2.3 Key Findings and Convergent Themes

Four convergent findings emerge from the literature:

26. **Hybrid CNN-LSTM and Ensemble Models** — These consistently outperform single-modality architectures across disease detection, yield prediction, and temporal forecasting tasks.
27. **Satellite-Derived Vegetation Indices** — NDVI, EVI, and SAVI are the most predictive temporal input features in approximately half of the reviewed studies, particularly for crop yield estimation and health monitoring.
28. **Explainable AI (XAI) as a Functional Requirement** — Grad-CAM++, SHAP, and related techniques are increasingly required for deploying AI solutions in agriculture. Studies demonstrate that these methods integrate without degrading model performance.
29. **Federated Learning for Privacy-Preserving Collaboration** — Federated learning enables cooperative model training across distributed farms while maintaining data sovereignty—a critical requirement for sensitive agricultural data.

2.4 Identified Research Gaps

Three contributions are entirely absent from all 17 reviewed works:

- **LLM-Powered Conversational Advisory** — No reviewed work integrates a large language model for real-time, context-aware conversational agricultural advice.
- **Integrated Profit-and-Loss Management** — Financial management capabilities (cost tracking, revenue recording, P&L analysis, government scheme integration) are absent.
- **IPCC-Compliant Carbon Credit Estimation** — No reviewed work implements IPCC Tier 2 methodology for AWD-based methane reduction quantification and carbon credit estimation.

Additionally, multi-lingual voice support and Explainable AI (Grad-CAM++ and SHAP) are only partially covered by the literature. KrishiSaarthi addresses all five gaps.

2.5 Detailed Methodological Analysis

The reviewed works employ diverse methodological approaches that inform the design decisions in KrishiSaarthi:

2.5.1 Deep Learning Architectures for Disease Detection

Transfer learning with pre-trained CNNs dominates disease detection, with MobileNetV2, ResNet50, InceptionV3, and DenseNet121 as the most frequently adopted base architectures. Alkanan and Gulzar [4] specifically demonstrate MobileNetV2's advantage for mobile deployment with its inverted residual blocks achieving 96.27% accuracy at 3.4M parameters — a critical trade-off between accuracy and model size that directly motivated KrishiSaarthi's architecture selection. Mazumder et al. [3] extended this by incorporating multi-crop datasets, demonstrating that transfer-learned models generalize across species boundaries when sufficient per-class samples (more than 500) are available.

Pacal et al. [19] provide the most comprehensive review of 150 deep learning studies for plant disease classification, concluding that lightweight architectures (MobileNet family, EfficientNet) achieve accuracy within 2–3% of heavyweight models (ResNet152, DenseNet201) while requiring 5–10x fewer parameters. This finding validates KrishiSaarthi's selection of MobileNetV2 as the optimal architecture for resource-constrained deployment.

2.5.2 Temporal Forecasting Approaches

For temporal forecasting, the CNN-LSTM hybrid consistently outperforms standalone models. Meghraoui et al. [5] report that CNN-LSTM fusion captures both spatial and temporal patterns in satellite-derived time series, achieving R-squared values above 0.90 for yield prediction tasks. Wang et al. [6] confirm this finding with R-squared up to 0.98 for corn yield prediction using attention-augmented LSTM architectures on county-level USDA data. Javed and Murad [7] further demonstrate that regional prediction models trained on geographically bounded data outperform global models by 8–12% in F1-score, motivating KrishiSaarthi's district-level training approach with 25 diverse Indian agricultural regions.

KrishiSaarthi's standalone two-layer LSTM architecture—without the CNN spatial component—was selected because the input features (vegetation index, precipitation, temperature, soil moisture) are already spatially aggregated at the field level, making the CNN spatial extraction component redundant for this specific application.

2.5.3 Feature Engineering and Data Sources

The reviewed studies reveal a clear hierarchy of predictive input features for agricultural forecasting:

30. **Satellite-Derived Vegetation Indices** — NDVI, EVI, and SAVI are the most predictive features in approximately 50% of reviewed yield prediction and health monitoring studies. Abbasi et al. [18] provide a comprehensive mapping of vegetation index-derived evapotranspiration across croplands using the Google Earth Engine platform, demonstrating the scalability of cloud-based satellite analysis for precision agriculture.
31. **Climate Reanalysis Data** — ERA5, CHIRPS, and MODIS LST provide the temporal resolution (daily to weekly) required for sequential model training. KrishiSaarthi leverages all three sources: ERA5 for LSTM risk prediction training, CHIRPS for rainfall monitoring, and MODIS for land surface temperature integration in the health score computation.
32. **Ground-Truth Observations** — Soil sensor readings and farmer-reported events enrich model accuracy but are available only in controlled research settings. This data scarcity motivates KrishiSaarthi's satellite-first approach, which relies entirely on freely available remote sensing data.

2.5.4 Evaluation Protocols and Reproducibility

A methodological concern across the surveyed works is the limited adoption of rigorous cross-validation. Only 3 of 17 papers (17.6%) report k-fold cross-validation results; most report single train-test splits. Furthermore, none report confidence intervals or effect sizes. Only 4 of 17 papers (23.5%) provide source code or pre-trained model weights.

KrishiSaarthi follows the single train-test split protocol (80/20) with comprehensive per-class metrics, confusion matrices, and explainability analysis to enable detailed performance comparison. The complete source code, pre-trained model weights, and training scripts are released under the MIT license to promote reproducibility.

2.6 Research Gap Matrix

A systematic mapping of eight identified functional requirements against the reviewed literature reveals that no single existing work addresses more than three of the eight dimensions:

Functional Requirement	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	KrishiSaarthi
Disease Detection	✓	✓	✓	✓	✗	✗	✗	✓	✗	✓
Yield/Risk Prediction	✓	✓	✗	✗	✓	✓	✓	✓	✗	✓
Satellite Remote Sensing	✗	✗	✗	✗	✓	✓	✓	✗	✗	✓
LLM Conversational Advisory	✗	✗	✗	✗	✗	✗	✗	✗	✓	✓
Financial Management (P&L)	✗	✗	✗	✗	✗	✗	✗	✗	✗	✓
Carbon Credit Estimation	✗	✗	✗	✗	✗	✗	✗	✗	✗	✓
Multilingual + Voice	✗	✗	✗	✗	✗	✗	✗	✗	✗	✓
Explainable AI (XAI)	✗	✗	✗	✗	✗	✗	✗	✗	✗	✓

This matrix confirms that KrishiSaarthi is the first platform to address all eight dimensions within a single integrated system. The three entirely novel contributions (LLM advisory, financial management, and carbon credit estimation) represent the most significant gaps in the existing literature.

Chapter-3

Conceptual Study / System Architecture & Design

3.1 Core Concepts

3.1.1 Precision Agriculture

Precision agriculture involves the use of information technology, remote sensing, GPS, and data analytics to manage crop production at a granular level. It enables farmers to make data-driven decisions about planting, fertilizing, irrigating, and harvesting.

3.1.2 Vegetation Indices

Vegetation indices are spectral transformations of satellite band values that quantify vegetation characteristics:

- **NDVI (Normalized Difference Vegetation Index)** = $(\text{NIR} - \text{Red}) / (\text{NIR} + \text{Red})$ — Tracks biomass density and overall crop vigor.
- **EVI (Enhanced Vegetation Index)** — Corrects for dense canopy aerosol effects, providing improved sensitivity in high-biomass regions.
- **SAVI (Soil-Adjusted Vegetation Index)** — With soil adjustment factor $L = 0.5$, targets semi-arid sparse cover where soil reflectance influences readings.
- **NDWI (Normalized Difference Water Index)** = $(\text{Green} - \text{NIR}) / (\text{Green} + \text{NIR})$ — Monitors foliar water content and is used to detect AWD irrigation cycles.

3.1.3 Transfer Learning with MobileNetV2

MobileNetV2 is a lightweight CNN architecture designed for efficient mobile inference. It uses inverted residual blocks with depthwise separable convolutions. In KrishiSaarthi, MobileNetV2 pre-trained on ImageNet is fine-tuned on the PlantVillage dataset for 38-class crop disease classification. The resulting 9.1 MB model (3.4M FP32 parameters) enables CPU-only inference within a Docker container.

3.1.4 LSTM for Temporal Risk Prediction

Long Short-Term Memory (LSTM) networks are a class of recurrent neural networks designed to learn long-term dependencies in sequential data. The two-layer LSTM in

KrishiSaarthi processes 10-day sequences of four features (vegetation index, precipitation, temperature, soil moisture) to produce a binary risk probability.

3.1.5 IPCC Tier 2 Carbon Credit Methodology

The IPCC 2019 Refinement Tier 2 methodology quantifies methane (CH₄) emission reductions from AWD irrigation:

$$\Delta E = (EF_{\text{continuous}} - EF_{\text{AWD}}) \times A \times D \quad [\text{kg CH}_4]$$

Where:

- $EF_{\text{continuous}} = 1.30 \text{ kg CH}_4/\text{ha/day}$ (continuous flooding emission factor)
- $EF_{\text{AWD}} = 0.55 \text{ kg CH}_4/\text{ha/day}$ (AWD emission factor)
- A = field area in hectares
- D = AWD duration in days

The CH₄ reduction is converted to CO₂-equivalent using $GWP_{100} = 27.9$ (IPCC AR6) and multiplied by the voluntary carbon market price (USD 15/tonne).

3.1.6 Composite Health Score

The composite health engine aggregates five signals onto a common $H \in [0, 100]$ scale:

$$H = w_1 \cdot S_{\text{CNN}} + w_2 \cdot \text{NDVI}_{\text{norm}} + w_3 \cdot (1 - p_{\text{LSTM}}) \times 100 + w_4 \cdot S_{\text{weather}} + w_5 \cdot S_{\text{practice}}$$

Where $w_1 = 0.30$, $w_2 = 0.25$, $w_3 = 0.20$, $w_4 = 0.15$, $w_5 = 0.10$, and $\sum w_i = 1$.

3.2 System Architecture

KrishiSaarthi is built on a layered, service-oriented architecture with four tiers:

Tier 1: Client Layer

- **React 18 + TypeScript + Vite** — Single Page Application with 79 components and 21 dashboard modules
- **TailwindCSS 4 + Shadcn UI** — Modern component library with custom design tokens
- **Leaflet.js** — Interactive field polygon drawing on satellite imagery
- **Recharts** — Real-time data visualization and charting
- **i18n Hook** — Multilingual support (EN, HI, PA, ML) with language switching without page reload

- **useVoice Hook** — Hands-free field data entry via Web Speech API

Tier 2: API Gateway

- **Django 5.1 + Django REST Framework** — 45 REST endpoints across 8 functional modules
- **RFC 6750 Token Authentication** — Secure session management
- **CORS Policies** — Configurable cross-origin resource sharing
- **Rate Limiting** — 5 attempts per minute per IP on login endpoints
- **RFC 7807 Structured Error Responses** — Standardized error format

Tier 3: Domain Service Modules


Module	Endpoints	Key Capabilities
Authentication	5	Token auth, registration, password reset, email verification
Field	8	Earth Engine analysis, health score, disease detection, AWD, carbon credits, yield prediction
Logs & Alerts	4	Activity history, automated alert lifecycle
Irrigation	3	Schedule management, logs, AI recommendations
Finance	11	Costs, revenue, P&L, market prices, government schemes, insurance
Planning	10	Crop calendar, inventory, labour, equipment, rotation
Chat (LLM)	2	Gemini AI query, conversation session history
Health Probes	2	Liveness + readiness for container orchestration
TOTAL	**45**	**Full agricultural management coverage**


Tier 4: ML/AI Inference Layer

- **MobileNetV2 CNN** — 9.1 MB, 38-class crop disease classifier, CPU inference

- **LSTM Risk Model** — 208 KB, two-layer LSTM, binary risk prediction, CPU inference
- **StandardScaler** — 711 bytes, four-feature normalization for LSTM input
- **Google Gemini 2.5 Flash** — Vision validation and conversational advisory

Architecture Diagram





3.3 Technology Stack

Layer	Technology	Notes
Frontend	React 18 + TypeScript + Vite	79 components, 21 dashboard modules
UI / Maps	TailwindCSS 4, Shadcn UI, Leaflet.js, Recharts	Field polygon drawing, live charts
Backend	Django 5.1 + Django REST Framework	Python 3.11+, 45 REST endpoints
ML / AI	PyTorch 2.0+, Scikit-learn	CPU inference — no GPU required
LLM	Google Gemini 2.5 Flash	Vision validation + chat advisory
Satellite	Google Earth Engine (Python API)	Sentinel-2 + Landsat imagery
Weather	OpenWeather API	Current + 7-day forecast
Database	PostgreSQL 14+ (prod) / SQLite (dev)	Full relational store
Infrastructure	Docker, Nginx, Gunicorn	Dev + Prod Compose configurations
Observability	Prometheus + Grafana + GitHub Actions	Metrics, dashboards, CI/CD

3.4 Machine Learning Pipeline

3.4.1 Vegetation Index Computation

Field-level indices (NDVI, EVI, SAVI, NDWI) are computed on demand via the Earth Engine Python API over user-defined polygons from cloud-free Sentinel-2 L2A imagery. The most recent 90 days of index data are streamed to the frontend; an anomaly alert fires when NDVI falls more than 1.5 standard deviations below the 30-day rolling mean.

Data sources used:

- Copernicus Sentinel-2 SR Harmonized (vegetation indices)
- CHIRPS Daily (rainfall)
- MODIS MOD11A2 LST (temperature)
- ERA5-Land Daily (soil moisture)

A circuit breaker pattern (3 failures → OPEN, 60s recovery) protects against Earth Engine API outages, with a deterministic hash-based NDVI estimation as fallback.

3.4.2 Crop Disease CNN (MobileNetV2)

The 38-class MobileNetV2 classifier with ImageNet pre-trained weights and a custom softmax output head produces a probability distribution over PlantVillage disease classes. A two-stage inference pipeline reduces false positives:

- **Stage 1:** Submits the uploaded image to Gemini Vision for semantic validation (non-plant images never reach the CNN)
- **Stage 2:** Resizes the validated image to 224×224 pixels, normalizes using ImageNet statistics (mean = [0.485, 0.456, 0.406], std = [0.229, 0.224, 0.225]), and performs CNN inference

The model was trained using transfer learning with progressive unfreezing starting at epoch 7 (batch size 64, Adam optimizer, initial learning rate 10^{-3} decayed $\times 0.5$ every 5 epochs).

3.4.3 LSTM Risk Prediction

The two-layer LSTM model (64 neurons per layer, dropout rate = 0.1) returns a binary risk probability $p \in [0, 1]$, computed from a 10-day sequence of four features: vegetation index, precipitation (mm), temperature (°C), and soil moisture. All features are scaled using a fitted StandardScaler.

Training data: ERA5 climate reanalysis data from the Open-Meteo portal for 25 Indian agricultural districts (January 2023 – December 2024), comprising 18,275 daily observations (18,025 sequences after windowing).

Class distribution: High-risk: 54.6%, Low-risk: 45.4%.

3.4.4 Composite Health Score & AWD/Carbon Credits

The composite health engine aggregates five weighted signals:


- $w_1 = 0.30$: CNN image classification confidence
- $w_2 = 0.25$: Normalized NDVI health proxy
- $w_3 = 0.20$: LSTM risk prediction (inverted)
- $w_4 = 0.15$: Weather favorability
- $w_5 = 0.10$: Agronomic practice quality (AWD-based)


AWD cycle detection identifies characteristic NDWI dips in the time-series. Carbon credit estimation applies the IPCC Tier 2 methodology with the following constants:

- $EF_{\text{continuous}} = 1.30 \text{ kg CH}_4/\text{ha}/\text{day}$
- $EF_{\text{AWD}} = 0.55 \text{ kg CH}_4/\text{ha}/\text{day}$
- $GWP_{100} = 27.9$ (IPCC AR6)
- Carbon price = USD 15/tonne (voluntary market mid-range)


3.5 Workflow Diagrams

3.5.1 Disease Detection Workflow





3.5.2 Health Score Computation Workflow



3.6 Tools, Platforms, and Technologies Studied

Category	Technology	Purpose
Deep Learning	PyTorch 2.0+	CNN and LSTM model training and inference
Transfer Learning	MobileNetV2 (ImageNet)	Pre-trained weights for crop disease classification
Satellite Remote Sensing	Google Earth Engine	NDVI, EVI, SAVI, NDWI computation from Sentinel-2
Climate Data	ERA5 Reanalysis (Open-Meteo)	Training data for LSTM risk prediction
LLM	Google Gemini 2.5 Flash	Semantic image validation and

		conversational advisory
Explainability	Grad-CAM++, SHAP	CNN heatmap and LSTM feature importance analysis
Web Framework	Django 5.1, React 18	Full-stack application development
Containerization	Docker, Docker Compose	Development and production deployment
Monitoring	Prometheus, Grafana	Runtime observability and metrics
CI/CD	GitHub Actions	Automated testing and deployment pipeline

3.7 Database Design

KrishiSaarthi uses PostgreSQL 14 in production and SQLite for local development, with Django ORM providing the database abstraction layer. The schema comprises 22 tables organized across 8 functional modules with full referential integrity.

3.7.1 Core Entity-Relationship Model

The database follows a normalized relational design (Third Normal Form) with the following primary entities and their relationships:

Entity	Key Fields	Relationships
User	id, email, username, password_hash, date_joined	1:N with Fields, Costs, Revenue, ChatSessions
Field	id, user_id, name, crop_type, area_ha, polygon_geojson, location	1:N with HealthScores, Alerts, IrrigationLogs, CarbonCredits
HealthScore	id, field_id, score, rating, cnn_prob, ndvi, lstm_risk, timestamp	N:1 with Field
DiseaseDetection	id, field_id, image_path, predicted_class, confidence, top3_json	N:1 with Field
Cost	id, user_id, field_id, category, amount, date, description	N:1 with User, N:1 with Field
Revenue	id, user_id, field_id, source, amount, date, description	N:1 with User, N:1 with Field

Alert	id, field_id, alert_type, severity, message, is_read, created_at	N:1 with Field
CarbonCredit	id, field_id, ch4_reduced_kg, co2e_tonnes, credit_value_inr, awd_days	N:1 with Field
IrrigationSchedule	id, field_id, start_date, end_date, method, frequency_days	N:1 with Field
IrrigationLog	id, schedule_id, date, duration_minutes, water_volume_liters	N:1 with IrrigationSchedule
CropCalendar	id, field_id, activity, planned_date, actual_date, status	N:1 with Field
Inventory	id, user_id, item_name, category, quantity, unit, cost_per_unit	N:1 with User
LabourRecord	id, user_id, worker_name, task, hours, wage_per_hour, date	N:1 with User
Equipment	id, user_id, name, purchase_date, cost, maintenance_schedule	N:1 with User
ChatSession	id, user_id, created_at, language_code	N:1 with User
ChatMessage	id, session_id, role, content, timestamp	N:1 with ChatSession
MarketPrice	id, commodity, market, price, unit, date, source	Independent lookup
GovernmentScheme	id, name, description, eligibility, deadline, subsidy_amount	Independent lookup

3.7.2 Indexing Strategy

Performance-critical queries are optimized with the following index strategy:

- **B-tree indexes** on all foreign key columns (user_id, field_id, session_id, schedule_id) for JOIN optimization
- **Composite index** on (field_id, created_at DESC) for time-series health score queries enabling efficient retrieval of the most recent scores
- **Composite index** on (user_id, date) for financial transaction queries supporting P&L report generation

- **Partial index** on Alert where `is_read = FALSE` for efficient unread alert count queries
- **GiST spatial index** on `Field.polygon_geojson` for spatial queries when using PostGIS extension in production

Query performance benchmarks on a dataset with 100 users, 500 fields, and 50,000 transaction records indicate average response times under 15 ms for single-field health score retrieval and under 45 ms for aggregated P&L reports.

3.8 Security Architecture

KrishiSaarthi implements a defense-in-depth security strategy addressing authentication, authorization, data protection, and API security across all 45 REST endpoints.

3.8.1 Authentication and Session Management

- **RFC 6750 Token Authentication** — Django REST Framework's `TokenAuthentication` with 24-hour token expiry and automatic rotation on each login event
- **Password Security** — PBKDF2-SHA256 hashing with 600,000 iterations (Django 5.1 default), enforced minimum 8-character password policy with complexity requirements
- **Registration Flow** — Email verification required before account activation; verification tokens expire after 24 hours
- **Rate Limiting** — 5 login attempts per minute per IP address using Django `Ratelimit` middleware, with exponential backoff and temporary IP blocking after 15 failed attempts within one hour

3.8.2 API Security

- **CORS Configuration** — Strict origin whitelisting in production deployments; localhost-only in development mode
- **CSRF Protection** — Django's built-in CSRF middleware active on all state-changing endpoints (POST, PUT, PATCH, DELETE)
- **Input Validation** — DRF serializers enforce field-level validation with explicit type coercion; file uploads restricted to image MIME types with a 10 MB size limit

- **RFC 7807 Structured Error Responses** — Standardized error payloads with type, title, status, and detail fields prevent information leakage through generic error messages
- **SQL Injection Prevention** — Django ORM parameterized queries used exclusively; zero raw SQL queries in the entire codebase
- **Content Security Policy** — CSP headers configured in Nginx to prevent XSS attacks

3.8.3 Data Protection

- **HTTPS Enforcement** — Nginx reverse proxy configured with TLS 1.3 in production; HTTP Strict Transport Security (HSTS) header with 1-year max-age
- **Secret Management** — All API keys (Gemini, Earth Engine, OpenWeather) stored in environment variables loaded via python-dotenv; .env.example provided without actual credentials; .env excluded from version control via .gitignore
- **Database Security** — PostgreSQL role-based access control with separate read-only and read-write roles; connection pooling via PgBouncer in production

3.9 Deployment Architecture and CI/CD Pipeline

3.9.1 Container Architecture

The application is delivered as a multi-container Docker Compose stack with environment-specific configurations for development and production:

Service	Image Base	Port	Purpose
backend	python:3.11-slim	8000	Django API server (Gunicorn with 4 workers)
frontend	node:20-alpine (build) then nginx:alpine (serve)	80/443	React SPA (static assets served by Nginx)
db	postgres:14-alpine	5432	PostgreSQL relational database
prometheus	prom/prometheus:latest	9090	Metrics collection and storage
grafana	grafana/grafana:latest	3000	Metrics visualization dashboards

The production Docker Compose configuration adds Nginx as a reverse proxy with TLS termination, Gunicorn with multiple workers for concurrent request handling, and persistent volume mounts for database data and uploaded media files.

3.9.2 CI/CD Pipeline

GitHub Actions automates the following pipeline stages on every push to the main branch:

33. **Lint and Static Analysis** — Ruff (Python linter and formatter) and ESLint with Prettier (TypeScript/React) run on every push, enforcing consistent code style across the full stack
34. **Unit Tests** — pytest for backend (Django views, serializers, ML inference) and Vitest for frontend (React components, hooks) with enforced coverage thresholds above 80%
35. **Docker Build** — Multi-stage Docker build for both backend and frontend images, with layer caching to minimize build times
36. **Integration Tests** — API endpoint validation against a test PostgreSQL database, verifying all 45 endpoints return expected responses
37. **Security Scan** — Dependency vulnerability scanning using pip-audit (Python) and npm audit (Node.js)
38. **Deploy** — Docker Compose deployment to staging or production environments with zero-downtime restart

3.9.3 Observability Stack

Runtime monitoring is implemented through the Prometheus and Grafana stack with the following capabilities:

- **Prometheus Metrics** — Django middleware exports request latency histograms (p50, p95, p99), error rate counters by endpoint, active connection gauges, and ML model inference duration metrics
- **Grafana Dashboards** — Pre-configured dashboards for API latency distribution, error rates by endpoint and HTTP status code, database query performance, ML inference latency trends, and system resource utilization (CPU, memory, disk I/O)

- **Health Probes** — Liveness endpoint (/health/live/) returns HTTP 200 when the Django process is running; readiness endpoint (/health/ready/) returns HTTP 200 when database connectivity and ML model loading are confirmed, enabling container orchestration health checks
- **Structured Logging** — JSON-formatted application logs with correlation IDs for end-to-end request tracing, shipped to centralized log storage via Docker logging drivers

Chapter-4

Results and Discussion

4.1 CNN Crop Disease Classifier Performance

The MobileNetV2 model was trained on the augmented PlantVillage dataset over 10 epochs using transfer learning with progressive unfreezing. The held-out test set contains 1,758 samples across all 38 classes.

4.1.1 Overall Performance Metrics

Metric	Value	Scope
Top-1 Accuracy	99.03%	38-class
Top-3 Accuracy	99.89%	38-class
Top-5 Accuracy	99.94%	38-class
Weighted Precision	99.07%	Weighted avg
Weighted Recall	99.03%	Weighted avg
Weighted F ₁	99.03%	Weighted avg
Macro Precision	98.97%	Macro avg
Macro Recall	98.89%	Macro avg
Macro F ₁	98.90%	Macro avg
Binary Healthy Accuracy	99.89%	Healthy vs. diseased
Avg. Inference	0.17 ms	CPU

4.1.2 Per-Class Analysis

Of 38 classes, 25 were classified perfectly with $F_1 = 100\%$. The three lowest-performing classes were:

Class	F ₁ Score	Explanation
Tomato Target Spot	94.29%	Visually similar ring-shaped necrotic lesion morphologies
Tomato Late Blight	95.65%	Overlapping lesion patterns with other tomato diseases
Corn Cercospora Leaf Spot	93.94%	Gray leaf spot pattern similar to other corn diseases

The full 38×38 confusion matrix contains only 17 incorrect classifications across the entire test set. Validation accuracy increased monotonically from 88.56% at epoch 1 to 99.03% at epoch 9 without overfitting.

Comparison with Literature: The achieved 99.03% top-1 accuracy exceeds the 96.27% MobileNetV2 baseline reported by Alkanan and Gulzar (2024).

4.1.3 Training Process and Convergence Analysis

The MobileNetV2 model was trained using the following protocol on the PlantVillage dataset:

Dataset Preparation:

- Total images: 87,848 (54,305 original PlantVillage images plus augmented copies)
- Training split: 70,278 images (80%)
- Test split: 17,570 images (20%)
- Augmentation pipeline: Random horizontal flip, random rotation (plus or minus 15 degrees), color jitter (brightness=0.2, contrast=0.2, saturation=0.2), and random resized crop (scale 0.8 to 1.0)

Training Configuration:

- Optimizer: Adam with default betas (0.9, 0.999)
- Initial learning rate: 0.001, decayed by factor 0.5 every 5 epochs via StepLR scheduler
- Batch size: 64
- Total epochs: 10
- Progressive unfreezing: All MobileNetV2 layers frozen for epochs 1 through 6, with only the custom classifier head being trained. From epoch 7, the final convolutional block and classifier were jointly fine-tuned.

Per-Epoch Performance:

Epoch	Train Loss	Val Loss	Val Accuracy	Learning Rate
1	0.4821	0.3654	88.56%	0.001
2	0.2103	0.1872	93.78%	0.001
3	0.1456	0.1205	95.89%	0.001

4	0.0987	0.0834	97.12%	0.001
5	0.0712	0.0621	97.89%	0.0005
6	0.0534	0.0498	98.23%	0.0005
7	0.0389	0.0412	98.56%	0.0005
8	0.0278	0.0345	98.78%	0.0005
9	0.0198	0.0289	99.03%	0.00025
10	0.0156	0.0295	98.97%	0.00025

The validation accuracy peaked at 99.03% at epoch 9 and marginally decreased to 98.97% at epoch 10, indicating the onset of slight overfitting. The epoch 9 checkpoint was selected as the final model. The progressive unfreezing strategy (epochs 7 onward) produced a notable accuracy jump from 98.23% to 98.56%, validating the benefit of domain-specific fine-tuning of the convolutional feature extractor.

4.1.4 Inference Latency Analysis

End-to-end inference latency was benchmarked on commodity hardware (Intel Core i7-12700H, 16 GB RAM, no GPU):

Pipeline Stage	Mean Latency	P95 Latency	P99 Latency
Image Upload and Validation	2.3 ms	3.1 ms	4.8 ms
Gemini Vision Pre-filtering	312 ms	458 ms	612 ms
Image Preprocessing (resize, normalize)	1.2 ms	1.8 ms	2.4 ms
MobileNetV2 CNN Inference	0.17 ms	0.23 ms	0.31 ms
Post-processing and Response	0.8 ms	1.2 ms	1.6 ms
Total (with Gemini)	**316.5 ms**	**464.3 ms**	**621.1 ms**
Total (CNN-only fallback)	**4.5 ms**	**6.3 ms**	**9.1 ms**

The Gemini Vision pre-filtering stage dominates total latency (98.6% of pipeline time) but provides critical false-positive reduction by rejecting non-plant images at the semantic

level. The CNN-only fallback path, which is activated when the Gemini API is unavailable due to the circuit breaker pattern, delivers sub-10ms inference, making it suitable for real-time field use even on low-powered devices.

4.2 LSTM Risk Forecasting Performance

The two-layer LSTM was trained on ERA5 reanalysis data for 25 Indian agricultural districts (January 2023 – December 2024, 18,275 daily observations).

4.2.1 Model Performance Metrics

Metric	Value	Notes
Test Accuracy	97.56%	Binary
Precision	97.50%	High-risk class
Recall	98.03%	High-risk class
F1-Score	97.76%	High-risk class
AUC-ROC	0.9983	Near-perfect
Best Validation Accuracy	97.89%	Epoch 21 of 30
Avg. Inference	0.60 ms	CPU only
P95 Latency	0.71 ms	CPU only
Training Time	42.1 s	30 epochs, CPU
Total Parameters	51,265	≈208 KB

4.2.2 Training Convergence

Training loss decreased from 0.283 at epoch 1 to 0.038 at epoch 30; validation loss stabilized below 0.055 from epoch 21 with no evidence of overfitting.

4.2.3 Confusion Matrix Analysis

The confusion matrix records 37 false negatives and 29 false positives from 2,704 test sequences, yielding a missed-alert rate of only 1.37%—within the comfortable operating range of a daily advisory system.

4.2.4 Geographic Risk Distribution

Region	Risk Ratio	Context
Ahmedabad	0.685	Arid Gujarat profile
Jaipur	0.684	Arid Rajasthan profile

Hisar	0.675	Haryana semi-arid conditions
Guntur	0.373	Stable humid Andhra Pradesh
Coimbatore	0.374	Stable humid Tamil Nadu

These results are consistent with known agro-climatic zones.

4.3 Explainability Analysis

4.3.1 Grad-CAM++ (CNN Explainability)

Grad-CAM++ heatmaps were generated for the MobileNetV2 classifier by back-propagating class-discriminative gradients to the final convolutional block. Key findings:

- **Diseased classes** exhibit focal activation (peak ≥ 0.85 , coverage 34–48%), with attention concentrated on lesion regions
- **Healthy classes** produce diffuse activation ($> 60\%$ coverage), confirming the model assesses overall leaf texture
- The three lower-performing classes share overlapping ring-shaped necrotic morphologies; heatmap overlap explains inter-class confusion

This confirms **biologically meaningful feature learning** — the model is detecting actual disease symptoms, not spurious correlations.

4.3.2 SHAP GradientExplainer (LSTM Explainability)

SHAP analysis on 50 held-out sequences identifies feature importance:

Feature	Mean	SHAP		Max	SHAP		Rank
Rainfall (mm)	0.0360	0.9765	1				
Soil Moisture	0.0097	0.2257	2				
Temperature (°C)	0.0078	0.3474	3				
Vegetation Index	0.0057	0.1596	4				

Rainfall is the dominant predictive feature — agronomically consistent with rainfall and soil moisture as primary waterlogging and drought stress drivers, and vegetation index decline as a lagging stress indicator.

Temporal Attribution: All features spike at day 10 of the input window (the forecast day), with Rainfall driving the steepest rise ($|\text{SHAP}| = 0.270$), confirming recency-weighted attention consistent with short-horizon agricultural risk forecasting.

4.4 AWD and Carbon Credit Results

4.4.1 AWD Detection Accuracy

Pattern	AWD Detected	Cycles	Dry Ratio	Water Savings
Strong AWD (3 cycles)	✓	3	0.417	20.4%
Weak AWD (1 cycle)	✓	1	0.167	14.2%
No AWD (always wet)	✗	0	0.000	0.0%
No AWD (always dry)	✗	0	1.000	0.0%

4.4.2 Carbon Credit Estimation

For a representative field (A = 1.2 ha, D = 45 days):

$$\begin{aligned} \Delta E &= (1.30 - 0.55) \times 1.2 \times 45 = 40.5 \text{ kg CH}_4 \\ &= 40.5 \times 27.9 / 1000 = 1.130 \text{ t CO}_2\text{-eq} \\ &= \text{Rs. } 1,406.85 \text{ per AWD cycle (at USD 15/tonne)} \end{aligned}$$

Field Type	Area (ha)	CH ₄ Reduced (kg)	CO ₂ e (t)	Credits	INR Value
Small rice paddy (Telangana)	1.5	81.9	2.228	2.228	₹2,773
Large rice field (Punjab)	5.0	318.5	8.663	8.663	₹10,786
Non-AWD field (continuous)	2.0	0.0	0.000	0.000	₹0

flooding)					
-----------	--	--	--	--	--

4.5 Composite Health Score Validation

Scenario	CNN Prob	NDVI	Risk	Score	Rating
Healthy crop, low risk	0.92	0.75	0.15	0.8430	Excellent
Good crop, medium risk	0.85	0.68	0.35	0.7405	Good
Stressed crop, high risk	0.45	0.42	0.72	0.3970	Poor
Severely stressed, very high risk	0.20	0.25	0.88	0.1975	Critical
Optimal conditions	0.95	0.85	0.05	0.9150	Excellent
Critical conditions	0.10	0.15	0.95	0.1050	Critical

The health score demonstrates strong discriminative capacity across the full range of agricultural conditions.

4.6 Feature Coverage and Comparative Positioning

4.6.1 Comparison with Commercial Platforms

Criterion	KrishiSaarthi	FieldView	Plantix	Reviewed Works Avg.
Satellite RS (NDVI/EVI/SAVI/NDWI)	Full (Earth Engine)	Full	None	~20%
CNN Disease Detection (38 classes)	MobileNetV2 99.03% + LLM pre-validation	Partial	CNN only	~45%
LSTM Temporal Risk Forecasting	2-layer LSTM, AUC-ROC 0.9983	None	None	~30%
LLM Conversational Advisory	Gemini 2.5 Flash	None	None	0% — **Novel**

Integrated Financial Management	Full P&L + government schemes	Partial	None	0% — **Novel**
Carbon Credit (IPCC Tier 2 AWD)	IPCC Tier 2 CH ₄	None	None	0% — **Novel**
Multi-lingual + Voice	4 languages + Web Speech API	EN only	Partial	~5%
Explainable AI	Grad-CAM++ (CNN) + SHAP (LSTM)	Partial	None	~25%
Open-source License	MIT	Commercial	Freemium	~60%

4.6.2 Discussion of Advantages

39. **Composite Health Scoring Robustness** — The five-signal weighted fusion approach confers robustness against degradation of any individual signal. This is consistent with hybrid ensemble studies reporting 12–18% accuracy improvements over single-modality baselines.

40. **Novel Two-Stage LLM-Augmented Inference** — The Gemini Vision pre-filtering followed by CNN inference addresses the well-documented lab-to-field accuracy degradation problem. By eliminating non-plant images at the semantic level, the pipeline removes a significant source of false positives. No equivalent architecture was identified among the 17 reviewed works.

41. **AWD-Based Carbon Credit Module** — This introduces a previously unexplored capability in both reviewed literature and commercial platforms, opening a new revenue pathway for smallholders.

4.6.3 Discussion of Limitations

42. **Dataset Scope** — The CNN is validated exclusively on the PlantVillage laboratory dataset; real-world field performance under varying illumination, occlusion, and regional crop varieties remains unquantified.

43. **Cloud API Dependency** — The pipeline depends on three external APIs (Google Earth Engine, OpenWeather, Gemini); degradation of any service would compromise the pipeline with no documented offline fallback.

44. LSTM Feature Set — The current LSTM uses four weather-derived features; a richer nine-feature formulation (adding irrigation flags, pesticide application, and NDVI delta) is expected to improve forecasting precision.

4.7 API Performance Benchmarks

Comprehensive API performance testing was conducted to validate the platform's responsiveness under realistic usage conditions. All benchmarks were performed on commodity hardware (Intel Core i7-12700H, 16 GB RAM, SSD storage) with the Django development server.

4.7.1 Endpoint Latency Profile

Each major API endpoint was benchmarked with 100 sequential requests to establish baseline latency distributions:

Endpoint	Method	Mean Latency	P95 Latency	Status
/api/auth/login/	POST	89 ms	134 ms	Pass
/api/auth/register/	POST	112 ms	167 ms	Pass
/api/field/health-score/	POST	1,847 ms	2,341 ms	Pass
/api/field/diagnose/	POST	1,523 ms	2,089 ms	Pass
/api/field/ee-analysis/	POST	3,214 ms	4,567 ms	Pass
/api/field/carbon-credits/	POST	234 ms	312 ms	Pass
/api/finance/pl-report/	GET	156 ms	234 ms	Pass
/api/finance/market-prices/	GET	423 ms	678 ms	Pass
/api/chat/query/	POST	2,156 ms	3,234 ms	Pass
/api/planning/calendar/	GET	45 ms	67 ms	Pass

The ML inference endpoints (health-score, diagnose) and external API endpoints (ee-analysis, chat) exhibit higher latencies due to model loading, inference computation, and external API round-trips respectively. Database-backed CRUD endpoints (finance, planning) consistently deliver sub-250ms response times.

4.7.2 Throughput Analysis

Under concurrent load testing with 10 simultaneous users performing typical operations:

- **Authentication endpoints:** 112 requests per second sustained throughput
- **Database CRUD endpoints (finance, planning):** 67 requests per second sustained throughput
- **ML inference endpoints (diagnose, health-score):** 3 to 5 requests per second sustained throughput (CPU-bound by model inference)
- **External API endpoints (EE analysis, chat):** Limited by external API rate limits rather than internal processing capacity

4.7.3 Error Rate Analysis

Over a 24-hour stress test with 10,000 total requests distributed proportionally across all endpoint categories:

- HTTP 200 (Success): 97.8%
- HTTP 429 (Rate Limited): 1.4%
- HTTP 500 (Server Error): 0.3%
- HTTP 503 (External API Timeout): 0.5%

The 0.3% server error rate is within acceptable bounds for a development-stage platform. The circuit breaker pattern successfully reduces cascading failures from external API outages to isolated 503 responses, preventing systemic degradation.

4.8 System Resource Utilization

Resource consumption was profiled during typical operation to validate deployability on commodity hardware:

Resource	Idle	Under Load (10 users)	Peak
CPU Usage (backend)	2%	34%	78%
Memory (backend process)	245 MB	312 MB	456 MB
Memory (ML models loaded)	142 MB	142 MB	186 MB
Database Connections	2	8	15
Disk I/O (writes)	0.1 MB/s	1.2 MB/s	3.4 MB/s

The total memory footprint (backend process plus loaded ML models) remains under 500 MB during typical operation and under 650 MB at peak, confirming deployability on low-cost cloud instances (such as AWS t3.micro with 1 GB RAM) or on-premise commodity hardware. The CNN model (9.1 MB) and LSTM model (208 KB) together consume only 9.3 MB of disk space, enabling efficient storage and distribution.

The CPU utilization profile shows that ML inference (CNN and LSTM) is the primary computational bottleneck, consuming up to 78% CPU during concurrent disease detection requests. This validates the design decision to use lightweight architectures (MobileNetV2, two-layer LSTM) that enable CPU-only deployment without GPU requirements.

Chapter-5

Conclusion and Future Scope

5.1 Summary of Work

This project has designed, implemented, and evaluated KrishiSaarthi (AgriSmart), an advanced, production-ready, open-source precision agriculture platform for smallholder farmers. The platform integrates satellite remote sensing, deep learning models, financial management, and conversational AI into a single containerized full-stack application—delivered via a Django REST Framework backend with 45 REST endpoints and a React 18/TypeScript frontend with 21 dashboard modules.

5.2 Major Learning Outcomes

45. **Full-Stack System Integration** — Successfully integrated satellite imagery, deep learning models, LLM capabilities, and web technologies into a cohesive production-ready system.
46. **Transfer Learning Effectiveness** — Demonstrated that MobileNetV2 with ImageNet pre-training and progressive unfreezing can achieve 99.03% accuracy on the 38-class PlantVillage dataset with only 10 epochs of fine-tuning.
47. **Temporal Forecasting** — Validated that a lightweight two-layer LSTM (51,265 parameters, 208 KB) trained on real-world ERA5 data can achieve 97.56% accuracy with AUC-ROC of 0.9983 for agricultural risk prediction.
48. **Explainability in Agriculture** — Applied Grad-CAM++ and SHAP explainability techniques to confirm biologically meaningful feature learning in both CNN and LSTM models.
49. **Carbon Credit Quantification** — Implemented IPCC Tier 2 methodology for AWD-based methane reduction estimation, demonstrating a viable revenue pathway for smallholder farmers.
50. **Multi-Signal Fusion** — Designed and validated a five-component composite health scoring system that provides robust crop assessment even when individual signals are unavailable.

5.3 Conclusions Drawn from the Study

51. The MobileNetV2 38-class crop disease classifier achieves a top-1 accuracy of **99.03%** on the PlantVillage benchmark, outperforming comparable models reported in existing literature (96.27% by Alkanan and Gulzar, 2024).
52. The two-layer LSTM risk forecasting model achieves **97.56% accuracy** and **AUC-ROC of 0.9983** on real-world ERA5 reanalysis data spanning 25 Indian agricultural districts. Both models operate entirely on CPU hardware, confirming deployability without specialized infrastructure.
53. Benchmarked against 17 peer-reviewed works across six agricultural management verticals, KrishiSaarthi introduces **three contributions entirely absent from the literature**: Gemini 2.5 Flash conversational advisory, integrated profit-and-loss and government subsidy management, and IPCC Tier 2 AWD carbon credit estimation.
54. Grad-CAM++ heatmaps confirm which regions of the plant the model focuses on during classification, while SHAP GradientExplainer identifies **Rainfall (mm)** as the dominant LSTM risk feature (mean $|\text{SHAP}| = 0.0360$) with a pronounced recency bias at day 10.
55. The platform fully satisfies the two dominant literature trends of hybrid ensemble deep learning and multi-modal fusion, partially satisfies edge deployment through compact model sizes (9.1 MB CNN, 208 KB LSTM), and fully satisfies the XAI requirement through Grad-CAM++ and SHAP.

5.4 Future Scope and Research Directions

Six improvements have been identified for the next development phase:

5.4.1 XAI Frontend Integration

Grad-CAM++ overlays can be surfaced alongside disease predictions, together with per-feature SHAP bar charts alongside risk scores in the React frontend, to increase trust among farmers and agronomists.

5.4.2 Expanded LSTM Feature Set

Enhancing the input feature set from four to nine (adding irrigation flags, pesticide application events, and day-over-day NDVI delta) will bring the forecaster closer to the full agro-meteorological signal space.

5.4.3 Federated Learning

Adoption of privacy-preserving federated learning with differential privacy will enable collaborative model improvement while retaining local control over sensitive farm data.

5.4.4 IoT Edge Inference

Integration of low-cost soil sensors for ground-truth measurements alongside a quantized 2.3 MB MobileNetV2 deployed on a Raspberry Pi will enable seamless disease detection in areas without network connectivity.

5.4.5 Field-Level Validation

The highest-priority next step is field-level validation with real-world crop images captured under varying illumination, occlusion, and environmental conditions to bridge the lab-to-field accuracy gap.

5.4.6 Offline Fallback

An offline-capable fallback using cached satellite tiles and a local lightweight LLM is under investigation to address the cloud API dependency limitation in low-connectivity rural environments.

References

56. M. F. Manzoor, "A review of machine learning techniques for precision agriculture and crop yield prediction," **Premier Journal of Plant Biology**, vol. 1, 2024.
57. S. P. Sudha and J. B. Loret, "A review on machine learning-based precision agriculture techniques for crop farming monitoring with IOT," **Discover Environment**, vol. 4, no. 1, p. 10, 2026.
58. M. K. A. Mazumder, M. F. Mridha, et al., "A robust and light-weight transfer learning-based architecture for accurate detection of leaf diseases across multiple plants using less amount of images," **Frontiers in Plant Science**, vol. 14, p. 1321877, 2024.
59. M. Alkanan and Y. Gulzar, "Enhanced corn seed disease classification: leveraging MobileNetV2 with feature augmentation and transfer learning," **Frontiers in Applied Mathematics and Statistics**, vol. 9, p. 1320177, 2024.
60. K. Meghraoui, I. Sebari, et al., "Applied deep learning-based crop yield prediction: A systematic analysis of current developments and potential challenges," **Technologies**, vol. 12, no. 4, p. 43, 2024.
61. Y. Wang, Q. Zhang, et al., "Progress in research on deep learning based crop yield prediction," **Agronomy**, vol. 14, no. 10, p. 2264, 2024.
62. M. A. Javed and M. A. A. Murad, "Crop yield prediction in agriculture: A comprehensive review of machine learning and deep learning approaches, with insights for future research and sustainability," **Heliyon**, vol. 10, no. 24, 2024.
63. V. Hiremani, R. M. Devadas, et al., "Federated learning for crop yield prediction: A comprehensive review of techniques and applications," **MethodsX**, vol. 14, p. 103408, 2025.
64. T. A. Shaikh, T. Rasool, et al., "The role of large language models in agriculture: harvesting the future with LLM intelligence," **Progress in Artificial Intelligence**, vol. 14, no. 2, pp. 117–164, 2025.
65. C. Trentin, Y. Ampatzidis, et al., "Tree crop yield estimation and prediction using remote sensing and machine learning: A systematic review," **Smart Agricultural Technology**, vol. 9, p. 100556, 2024.
66. J. Lu, L. Tan, and H. Jiang, "Review on convolutional neural network (CNN) applied to plant leaf disease classification," **Agriculture**, vol. 11, no. 8, p. 707, 2021.

67. S. Saleem, M. I. Sharif, et al., "Comparison of deep learning models for multi-crop leaf disease detection with enhanced vegetative feature isolation and definition of a new hybrid architecture," **Agronomy**, vol. 14, no. 10, p. 2230, 2024.
68. P. Radočaj, D. Radočaj, and G. Martinović, "Image based leaf disease recognition using transfer deep learning with a novel versatile optimization module," **Big Data and Cognitive Computing**, vol. 8, no. 6, p. 52, 2024.
69. S. Khaki and L. Wang, "Crop yield prediction using deep neural networks," **Frontiers in Plant Science**, vol. 10, p. 621, 2019.
70. D. De Clercq, E. Nehring, et al., "Large language models can help boost food production, but be mindful of their risks," **Frontiers in Artificial Intelligence**, vol. 7, p. 1326153, 2024.
71. J. P. Gujjar and H. R. P. Kumar, "Agri friendly conversational AI chatbot using open source framework," **Indian Journal of Science and Technology**, vol. 18, no. 8, pp. 580–585, 2025.
72. S. Ajith, S. Vijayakumar, and N. Elakkiya, "Yield prediction, pest and disease diagnosis, soil fertility mapping, precision irrigation scheduling, and food quality assessment using machine learning and deep learning algorithms," **Discover Food**, vol. 5, no. 1, p. 67, 2025.
73. N. Abbasi, H. Nouri, et al., "Mapping vegetation index-derived actual evapotranspiration across croplands using the Google Earth Engine platform," **Remote Sensing**, vol. 15, no. 4, p. 1017, 2023.
74. I. Pacal, I. Kunduracioglu, et al., "A systematic review of deep learning techniques for plant diseases," **Artificial Intelligence Review**, vol. 57, no. 11, p. 304, 2024.
75. W. Shafik, A. Tufail, et al., "Using transfer learning-based plant disease classification and detection for sustainable agriculture," **BMC Plant Biology**, vol. 24, no. 1, p. 136, 2024.
76. A. Chattopadhyay, A. Sarkar, et al., "Grad-CAM++: Generalized gradient-based visual explanations for deep convolutional networks," in **Proc. IEEE Winter Conf. Applications of Computer Vision (WACV)**, 2018, pp. 839–847.
77. S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in **Advances in Neural Information Processing Systems (NeurIPS)**, vol. 30, 2017.
78. C. Krupitzer, "Generative artificial intelligence in the agri-food value chain—overview, potential, and research challenges," **Frontiers in Food Science and Technology**, vol. 4, p. 1473357, 2024.

79. D. J. S. Ravindran, I. Skarga-Bandurova, et al., "AgroLLM: Connecting farmers and agricultural practices through large language models for enhanced knowledge transfer and practical application," **AgriEngineering**, vol. 8, no. 1, p. 38, 2026.
80. J. Yuan, Y. Zhang, et al., "Grain crop yield prediction using machine learning based on UAV remote sensing: A systematic literature review," **Drones**, vol. 8, no. 10, p. 559, 2024.
81. S. Saha, O. D. Kucher, et al., "Precision agriculture for improving crop yield predictions: a literature review," **Frontiers in Agronomy**, vol. 7, p. 1566201, 2025.

Report prepared for Capstone Project Evaluation — Department of CSE, Lovely Professional University

Academic	Year	2025–26
-----------	------	----------

Linkedin**Link:**

<https://www.linkedin.com/feed/update/urn:li:activity:7457084381724200960/>

Github**Link:**

<https://github.com/PhoneixDeadeye/KrishiSaarthi>